

1. (•) *I understand the course policies.*

Suppose you are given an array $A[1..n]$ of positive integers. Describe and analyze an algorithm to find the smallest positive integer that is not an element of A in $O(n)$ time.

Solution: A pseudo code of our algorithm is given as follows:

<pre>SMALLESTPOSITIVEINTEGER($A[1..n]$): for $i \leftarrow 1$ to $n + 1$ $B[i] \leftarrow 0$ for $i \leftarrow 1$ to n if $1 \leq A[i] \leq n$ $B[A[i]] \leftarrow 1$ for $i \leftarrow 1$ to $n + 1$ if $B[i] = 0$ return i</pre>
--

Running time: $O(n)$


Space: $O(n)$



2. (•) *I understand the course policies.*

Describe and analyze an efficient algorithm to find a solid block in M with maximum area.

Solution: We first observe that given a point at (i, j) , we can define a solid block by extending it until encountering a '0', then extend it to left right until the line reaches a '0'. If we can find all these solid blocks defined by (i, j) , then we can certainly find the largest solid block among them.



To do this, we first compute the largest solid blocks for a single row and column. Specifically, we define $top(i, j)$, $left(i, j)$ and $right(i, j)$ as the largest (unit width) block from (i, j) to the *top*, *left* and *right* side. For example, for a row $[1 \ 0 \ 1 \ 1 \ 1]$, its corresponding *left* array is $[1 \ 0 \ 1 \ 2 \ 3]$. We can use dynamic programming to compute it as follows:

$$left(i, j) = \begin{cases} left(i, j-1) + 1 & \text{if } M[i, j] = 1, \\ 0 & \text{otherwise.} \end{cases}$$

To handle the base case of $left(i, 1)$, we just add a sentinel column 0 and let $left(i, 0) = 0$. Similarly, $top(i, j)$ and $right(i, j)$ can be computed like $left(i, j)$. We use a 2D-array to store $top(i, j)$ (However, later we only need the *left* and *right* information for the row that we are working on, we can just use a 1D-array to store them). The evaluation order is just from *left* to *right* for $left(i, j)$, *right* to *left* for $right(i, j)$, and *top* down for $top(i, j)$.

Next, we compute a *left extension* and *right extension* for a given point (i, j) . They denote the *leftmost* and *rightmost* distance the block can extend to after reaching the *top*, respectively. Specifically, *left extension* can be simply computed as:

$$lext(i, j) = \min\{lext(i-1, j), left(i, j)\}$$

$rext(i, j)$ is defined similarly. For the base cases, we set $lext(0, j) = rext(0, j) = +\infty$ for all j as sentinel. Then, for each pair of i and j , we compute the area of solid block defined by (i, j) as $area(i, j) = (rext(i, j) - lext(i, j) + 1) \times top(i, j)$. We do not need to remember $area(i, j)$ at all because we can just keep track of the largest block we have ever found as we go. This takes $O(mn)$ time. After the scanning is over, we are done. The computation for $top(i, j)$, $left(i, j)$, $right(i, j)$, $lext(i, j)$ and $rext(i, j)$ can be finished in $O(mn)$ time. We can put them together to avoid excessive looping through (i, j) pairs. Hence, the overall runtime is $O(mn)$, so as the space. Note that for the $top[i, j]$, $lext[i, j]$ and $rext[i, j]$ array, since we only need the information for previous row, we can just use a sliding window technique to reduce the final space to $O(m)$. The pseudo code is described as follows. For clarity, we intentionally do not use the sliding window technique in it.

```

LARGESTSOLIDBLOCK( $M[1..m, 1..n]$ ):
  for  $j \leftarrow 1$  to  $n$             $\langle\langle$ Base cases $\rangle\rangle$ 
     $l_{ext}[0, j] \leftarrow +\infty$ 
     $r_{ext}[0, j] \leftarrow +\infty$ 
   $largest \leftarrow 0$ 
  for  $i \leftarrow 1$  to  $m$ 
     $left[0] \leftarrow 0$ 
    for  $j \leftarrow 1$  to  $n$         $\langle\langle$ Compute left array $\rangle\rangle$ 
      if  $M[i, j] = 1$ 
         $left[j] \leftarrow left[j - 1] + 1$ 
      else
         $left[j] \leftarrow 0$ 

     $right[n] \leftarrow 0$ 
    for  $j \leftarrow n$  down to 1  $\langle\langle$ Compute right array $\rangle\rangle$ 
      if  $M[i, j] = 1$ 
         $right[j] \leftarrow right[j + 1] + 1$ 
      else
         $right[j] \leftarrow 0$ 

    for  $j \leftarrow 1$  to  $n$ 
      if  $M[i, j] = 1$             $\langle\langle$ compute top array $\rangle\rangle$ 
         $top[i, j] \leftarrow top[i - 1, j] + 1$ 
      else
         $top[i, j] \leftarrow 0$ 
       $l_{ext}[i, j] \leftarrow \min\{l_{ext}[i - 1, j], left[j]\}$ 
       $r_{ext}[i, j] \leftarrow \min\{r_{ext}[i - 1, j], right[j]\}$ 
       $area \leftarrow (r_{ext}[i, j] - l_{ext}[i, j] + 1) \times top[i, j]$ 
       $largest \leftarrow \max\{largest, area\}$ 
  return  $largest$ 

```

■

3. (•) *I understand the course policies.*

Describe and analyze an algorithm to compute a maximum weight matching, given the tree T as input.

Solution: We assume that the problem statement is asking for a maximum matching in a weighted graph, i.e., match as many edges as possible, if there is tie, return the matching that has the largest weight. Given a tree T that is rooted in r . We denote it as $T(r)$. Let c be a child of r , we define $T(r) \setminus T(c)$ as the tree that is obtained by removing $T(c)$ from $T(r)$ as shown in Figure 1.





Figure 1. Illustration of $T(r)$ and $T(r) \setminus T(c)$.

Consider an edge $e = (u, r)$ in $T(r)$, there are two possible cases:

- (a) **If e is included in the maximum weight matching:** The maximum weight matching of T will then be the summation of the maximum weight matching of all the trees $T(x)$ for all x that is a child of u and the maximum weight matching of all the trees $T(y)$ for all y that is a child of r and $y \neq u$. An illustration is given in Figure 2(a).
- (b) **If e is not in the maximum weight matching:** The maximum weight matching of T will then be the summation of the maximum weight matching of $T(u)$ and the maximum weight matching of $T(r) \setminus T(u)$. An illustration is given in Figure 2(b).



(a) Edge e is included in the maximum weight matching.

(b) Edge e is not in the maximum weight matching.

Figure 2. Two cases for edge $e = (u, r)$

Hence, we can devise a dynamic programming algorithm to compute the maximum weight matching. We can use the tree itself to memoize the intermediate result. Specifically, for each node v , we can add a new field $v.MWM$ to store the weight of the maximum weight matching, and a new field $v.to$ to remember to which node v is matched. A pseudo code is given as follows:

```

MAXIMUMWEIGHTMATCHING( $T(v)$ ):
    if  $v$  has no child
        return 0
    if  $v.MWM$  is computed
        return  $v.MWM$ 
     $u \leftarrow$  any child of  $v$ 
     $e \leftarrow (u, v)$ 
     $withe \leftarrow w(e)$ 
    for each child  $w$  of  $u$ 
         $withe \leftarrow withe + \text{MAXIMUMWEIGHTMATCHING}(T(w))$ 
    for each child  $x$  of  $v$  and  $x \neq u$ 
         $withe \leftarrow withe + \text{MAXIMUMWEIGHTMATCHING}(T(x))$ 
     $withoute \leftarrow \text{MAXIMUMWEIGHTMATCHING}(T(u)) + \text{MAXIMUMWEIGHTMATCHING}(T(v) \setminus T(u))$ 
    if  $withe \geq withoute$ 
         $v.MWM = withe$ 
         $v.to = u$ 
    else
         $v.MWM = withoute$ 
    return  $v.MWM$ ;

```

Running time: Because each vertex is visited at most once, the running time of $\text{MAXIMUMWEIGHTMATCHING}(T(v))$ is $O(n)$. To output the maximum weight matching, we just perform a traversal to the tree T and output the matching pairs according to $v.to$ for all v . This also takes $O(n)$ time. Hence, the total running time is $O(n)$.

Space: Since we only add a constant amount of storage to each node and there are n nodes, the space complexity is $O(n)$. ■

4. (•) *I understand the course policies.*

Describe and analyze an algorithm that accepts three strings x , y , and z as input, and decides whether z is an interleaving of x and y .

Solution: We first *repeat* x to obtain a string X that has the same length as z , similarly we obtain a string Y by repeating y . We define $IsInterleave(X[1..i], Y[1..j], z[1..i+j])$ as follows:

$$IsInterleave(X[1..i], Y[1..j], z[1..i+j]) = \begin{cases} IsInterleave(X[1..i-1], Y[1..j], z[1..i+j-1]) & \text{if } X[i] = z[i+j], \text{ or} \\ IsInterleave(X[1..i], Y[1..j-1], z[1..i+j-1]) & \text{if } Y[j] = z[i+j] \\ false & \text{otherwise} \end{cases}$$


or, in a more compact form:

$$IsInterleave(X[1..i], Y[1..j], z[1..i+j]) = [X[i] = z[i+j]] IsInterleave(X[1..i-1], Y[1..j], z[1..i+j-1]) \vee [Y[j] = z[i+j]] IsInterleave(X[1..i], Y[1..j-1], z[1..i+j-1])$$

Let i and j denote the length contribution from x and y to z , respectively. In order to find if z is an interleaving of x and y , we only need to check all the possible combination of i and j , and returns *true* if any one of $IsInterleave(X[1..i], Y[1..j], z[1..n])$ is *true*. Clearly we can use a 2D-array to memoize the recurrence. This recurrence contains three base case:

- (a) It is trivially true when z is an empty string
- (b) It is true if $X[1..i]$ is a *prefix* of z for all $i \in [1..n]$ and $j = 0$
- (c) It is true if $Y[1..j]$ is a *prefix* of z for all $j \in [1..n]$ and $i = 0$

Hence, we can just slightly modify the original recurrence to $IsInterleave(X[0..i], Y[0..j], z[0..i+j])$, and add a row and a column with index 0 to the memoization data structure. The evaluation order is illustrated as follows:




Clearly, this algorithm costs $O(n^2)$ time and uses $O(n^2)$ space. ■

5. (•) **I understand the course policies.**


Describe and analyze an algorithm to compute the maximum total reward that the organizers could be forced to pay, given the array M as input.

Solution: We define the maximum reward we can get between a range of snails from i to j as $MaxReward(i, j)$, as the shaded region in the figure:



There are three possibilities for $MaxReward(i, j)$:


- (a) If snail i and j make up in the optimal solution, then the optimal solution is equal to $M[i, j] + MaxReward(i + 1, j - 1)$, where $M[i, j]$ is the reward for pair i and j .
- (b) If snail j is not pair up with some snail in range $[i..j]$, then the optimal solution is just equal to $MaxReward(i, j - 1)$.
- (c) If snail j pairs up with some snail in range $[i + 1..j - 1]$, then the optimal solution will be $\max\{MaxReward(i, k - 1) + MaxReward(k, j)\}$, for all k in $[i + 1..j - 1]$.



Hence, the maximum reward can be defined using following recursion:

$$MaxReward(i, j) = \max \left\{ \begin{array}{l} M[i, j] + MaxReward(i + 1, j - 1) \\ MaxReward(i, j - 1) \\ MaxReward(i, k) + MaxReward(k, j), \text{ for all } k \in [i + 1, j - 1] \end{array} \right\}$$

The optimal solution for this problem is $MaxReward(1, n)$. This recurrence has two simple base cases. When i is equal to j , since a single snail can not pair up with itself, thus we have $MaxReward(i, j) = 0$. Also, range $[i..j]$ is invalid when $i > j$, we also set them to 0. We can use a 2D-array to store the values. The evaluation order is illustrated as follows:



Clearly this algorithm will traverse all pair of i and j where $i < j$, during which a for loop is used to find the maximum value in the third case of the recurrence. Hence, it is a $O(n^3)$ algorithm, which uses $O(n^2)$ space. ■